



▶ Play Video

Memo

Backed by [Sequoia Capital](#), [Gladia](#) is a quickly growing AI audio infrastructure that transcribes and enriches every conversation through a single API—so developers can turn audio into structured, actionable data for their products.

Today, Gladia is the leading multilingual, real-time speech-to-text (STT) API. In the last two years, the company has processed **14M hours**, representing **nearly 1600 years of audio**. Gladia ended 2025 with ~95 enterprise subscriptions, over 2k self-serve customers, and **over 350k users globally**. The company hit **approx. \$4.5M ARR EO25 (2x+ YoY growth) → proj. \$10M+ ARR EO26 → proj. \$20M+ ARR EO27**.

We are investing on a **reasonable €52M valuation** alongside [XAnge](#) ([Shine](#), [Gleamer](#)), who is leading with a €1.5M check, [XTX Ventures](#) ([Groq](#), [Eko Health](#)), and [illuminate Financial](#) ([BeZero](#), [Genesis](#)), ahead of a planned \$25M Series B later this year. Existing investors include [Sequoia](#) ([Apple](#), [Instacart](#)), [GFC](#) ([Meta](#), [Revolut](#)), and [Soma Capital](#) ([Cruise](#), [Ramp](#)), among others.

Here is why we are excited to have secured allocation in Gladia...

Massive Market Opportunity: Gladia is targeting a large, underpenetrated speech AI market, driven by the proliferation of audio across customer calls, meetings, and emerging voice agent use cases. Less than 5% of global audio is currently processed by even basic speech-to-text (STT) systems, highlighting a massive greenfield for infrastructure providers as voice becomes a primary interface

transcription and audio intelligence layer underpinning next-generation conversational systems.

The Solution: [Gladia](#) is building the premier AI audio infrastructure for voice products. The company provides a developer-first speech AI infrastructure layer that delivers highly accurate, low-latency, real-time, and asynchronous transcription across 100+ languages, with native integrations into telephony, WebRTC, and modern voice agent frameworks. Its platform goes beyond basic STT by extracting structured data (e.g., emails, phone numbers), enabling scalable processing of thousands of concurrent audio streams, and meeting enterprise-grade requirements (SOC2, HIPAA, ISO27001). Built on a proprietary foundation model optimized for speed and multilingual performance, Gladia is designed to serve as the core audio data layer for call centers, meeting assistants, and AI voice agents.

Strong Team: Gladia is led by [Jean-Louis Quéguiner \(Co-founder & CEO\)](#), the former VP of AI at OVHcloud and CTO at Auchan. [Maxime Gaudin \(CPTO\)](#) is a former founder and experienced technical leader who previously helped scale Malt from €0 to €150M+ in GMV. [Thibaut Le Boulair \(COO\)](#)

Thibaut brings over a decade of experience managing complex industrial and operational projects at MBDA and Tallano Technologies. [Constant Dacbert \(Head of Sales\)](#) was previously Head of Sales at Crowdcube, an AE at Datadog, and a Sales Lead at Sscreen (acq. by Datadog).

Scaling Traction: Gladia is the leading multilingual, real-time speech-to-text API. In the last two years, the company has processed more than **14M hours**, representing **nearly 1600 years of audio**. Gladia ended 2025 with ~95 enterprise subscriptions, over 2k self-serve customers, and over **350k users globally**. The company hit **approx. \$4.5M ARR EO25 (2x+ YoY growth) → proj. \$10M+ ARR EO26 → proj. \$20M+ ARR EO27**.

Reasonable €52M Valuation + Notable Co-Investors: In the contrast of Gladia's ARR growth, we are investing on a **reasonable €52M valuation** alongside [XAnge \(Shine, Gleamer\)](#), who is leading with a €1.5M check, [XTX Ventures \(Grog, Eko Health\)](#), and [illuminate Financial \(BeZero, Genesis\)](#), ahead of a planned \$25M Series B later this year. Existing investors include [Sequoia \(Apple, Instacart\)](#), [GFC \(Meta, Revolut\)](#), and [Soma Capital \(Cruise, Ramp\)](#), among others.

[Invest in the MV RF for reduced syndicate carry & no cut-backs on oversubscribed syndicates.](#)

DEAL OVERVIEW

This is a €4M Series A extension on a €52M pre-money valuation ahead of a planned \$25M Series B later this year. Funds from this round will be used to:


- Scale to proj. \$10M+ ARR by EO26
- Make key U.S. commercial hires to accelerate GTM motion
- Advance R&D to secure long-term product differentiation

across customer calls, meetings, and emerging voice agent use cases. Less than 5% of global audio is currently processed by even basic speech-to-text (STT) systems, highlighting a massive greenfield for infrastructure providers as voice becomes a primary interface for AI applications. This gap, combined with the expansion of meeting assistants, call center automation, and autonomous voice agents, positions Gladia to capture value as a core real-time transcription and audio intelligence layer underpinning next-generation conversational systems.

COMPANY OVERVIEW

Gladia is a quickly growing AI audio infrastructure that transcribes and enriches every conversation through a single API—so developers can turn audio into structured, actionable data for their products.


Gladia provides a developer-first speech AI infrastructure layer that delivers highly accurate, low-latency, real-time, and asynchronous transcription across 100+ languages, with native integrations into telephony, WebRTC, and modern voice agent frameworks. Its platform goes beyond basic STT by extracting structured data (e.g., emails, phone numbers), enabling scalable processing of thousands of concurrent audio streams, and meeting enterprise-grade requirements (SOC2, HIPAA, ISO27001). Built on a proprietary foundation model optimized for speed and multilingual performance, Gladia is designed to serve as the core audio data layer for call centers, meeting assistants, and AI voice agents.



Gladia's foundation model, **Solaria**, is trained on **14M hours of real-world audio**, a dataset spanning **more than 100 languages**, providing both depth and diversity in linguistic representation. This large-scale, multilingual data foundation ensures that the company's models understand and adapt to global language nuances, dialects, and low-resource linguistic contexts, enabling superior performance across markets where traditional models often underperform. Solaria was designed with strong **accent resilience** and advanced capabilities, such as:

- **Automatic language detection**

This gives Gladia a significant edge in both **major markets** (English, French, German, Spanish, Italian, Portuguese) and **emerging ones** (Arabic, Asian languages), where demand is rapidly increasing.



Gladia's proprietary model development pipeline sets it apart:


- **Model architecture:** A large-scale **sequence-to-sequence architecture** designed for performance at an affordable cost.
- **Unique data curation:** In-house, high-quality data pipelines combining **automated large-scale unsupervised annotation** and **expert-guided curation**.
- **Training pipeline:** A proprietary stack inspired by **DeepSeek** and **Microsoft Phi-4**, engineered for efficient scaling on low-resource infrastructures.
- **Post-training innovation:** Advanced proprietary **Reinforcement Learning** and **Group Relative Policy Optimization (GRPO)** techniques that fine-tune model alignment and robustness.

Deep control over the full model lifecycle, data, training, and post-training, forms the backbone of Gladia's competitive IP and cost efficiency. As a result:

- The company delivers superior accuracy in the most popular languages.
- Gladia's real-time transcription is the fastest available, with an architecture optimised for scalability and low latency, **reaching sub-200ms**, which is lower than Deepgram.
- Gladia has already taken its latest model, Solaria-3 (yet to be released), further, **which is 30% more efficient** than Solaria-1 and **60% more cost-effective**.

Gladia's architecture also enables seamless integration with voice platforms, including Recall, LiveKit, Pipecat, and Vapi. These platforms are where modern voice products are built, from real-time agents to communication apps. By embedding Gladia directly within them, it can become the native voice intelligence layer powering their ecosystems with no custom setup.

Today, Gladia is the leading multilingual, real-time STT API. Tomorrow, it is evolving beyond a purely transactional API toward a relational speech intelligence platform, one that stays specialized in



GO TO MARKET STRATEGY


Gladia has a dual-motion growth model that combines sales-led execution (75% to 80% of total revenue) and a strong self-serve engine.

Gladia's self-serve growth relies on building trust and visibility within developer communities and technical SMBs / mid-market. This is accomplished through events, brand building, content marketing, and integrations with voice agent platforms such as Vapi, Jambonz, and automation workflow platforms such as Zapier and N8N to enable automation between transcription, metadata extraction, and downstream tools.

For large accounts, Gladia focuses on credibility, specialization, and partnerships. This includes recruiting recognized figures from voice AI, CCaaS, CPaaS, or BPO sectors as brand ambassadors, deepening partnerships with telephony and communication providers (e.g., Twilio, Aircall, Vonage, Telnyx) to position Gladia as the default speech-to-text provider, and embedding FDEs within partner engineering teams to identify prototype real-world solutions.

TRACTION & BUSINESS PERFORMANCE

- Hit **approx. \$4.5M ARR EO25 (2x+ YoY growth) → proj. \$10M+ ARR EO26 → proj. \$20M+ ARR EO27**
- The company ended 2025 with ~95 enterprise subscriptions, over 2k self-serve customers, and over 300k users globally
- Gladia has processed more than **14M hours** in about 2 years, representing **nearly 1600 years of audio**




Starter	Growth	Enterprise
Flexible pay-as-you-go for moderate audio volumes. Get started immediately.	Lower unit pricing for fast-growing teams. Commit upfront to unlock savings.	Annual plan with custom models, fine-tuning, debundled pricing, and more.
Async at \$0.61/hr Real-time at \$0.75/hr * 10h free monthly	Async as low as \$0.20/hr Real-time as low as \$0.25/hr * 67% less than Starter	Custom

COMPETITION

Today, Gladia is the leading multilingual, real-time speech-to-text API. Tomorrow, it is evolving beyond a purely transactional API toward a relational speech intelligence platform, one that stays specialized in speech AI, while going deeper into understanding human conversation at every level.

Gladia will become the audio data hub for every voice product.

Many competitors, including Deepgram, 11 Labs, Cartesia, and Hume, are moving toward voice-to-voice models that merge STT and TTS to serve the “Voice AI Generalist” market. This approach is highly transactional and risks competing with their own customers (voice-agent builders). Gladia focused on speech understanding, rather than voice agent creation. In doing so, it positions the company as the specialized speech layer powering agents instead of replacing them.



Gladia's strength lies in a deeply integrated approach that combines data excellence and proprietary model innovation. These pillars have enabled the company to build the best STT product on the market, uniquely adapted to real-world global communication, in two years and with five times fewer resources than direct competitors.


Over the past two years, Gladia has built one of the most comprehensive free-tier data ecosystems in the speech and text domains. The company's datasets span over 100 languages, providing both depth and diversity in linguistic representation. This large-scale multilingual data foundation ensures its models understand and adapt to global language nuances, dialects, and low-resource

Gladia's proprietary model development pipeline sets it apart:

- **Model architecture:** A large-scale **sequence-to-sequence architecture** designed for performance at an affordable cost.
- **Unique data curation:** In-house, high-quality data pipelines combining **automated large-scale unsupervised annotation** and **expert-guided curation**.
- **Training pipeline:** A proprietary stack inspired by **DeepSeek** and **Microsoft Phi-4**, engineered for efficient scaling on low-resource infrastructures.
- **Post-training innovation:** Advanced proprietary **Reinforcement Learning** and **Group Relative Policy Optimization (GRPO)** techniques that fine-tune model alignment and robustness.

This deep control over the full model lifecycle, data, training, and post-training, forms the backbone of Gladia's competitive IP and cost efficiency.

TEAM



Jean-Louis Quéguiner (Co-founder & CEO)

Jean-Louis has over a decade of experience in AI/ML. He previously served as VP of AI at OVHcloud, where he led large-scale AI initiatives, bringing deep expertise in building production-grade machine learning systems.

Maxime Gaudin (CPTO)

Maxime is a former founder and experienced technical leader who previously helped scale Malt from €0 to €150M+ in GMV. He specializes in building high-performance, scalable infrastructure and leading engineering teams through rapid growth.

Thibaut Le Boulair (COO)

Thibaut brings over a decade of experience managing complex industrial and operational projects. Before Gladia, he was a PM at MBDA and Tallano Technologies, where he developed expertise in

[Anna Jelezovskaia \(Head of Marketing\)](#)

Anna was previously a Content Strategist at Powder. She focuses on scaling brand, demand generation, and go-to-market strategies for developer-first and AI-native products.

[Constant Dacbert \(Head of Sales\)](#)

Constant was previously Head of Sales at Crowdcube, an AE at Datadog, and a Sales Lead at Sgreen (acq. by Datadog). He brings a strong background in enterprise SaaS sales and scaling revenue through high-performing GTM teams.

[Invest in the MV RF for reduced syndicate carry & no cut-backs on oversubscribed syndicates.](#)

Co-investors

	XAnge Lead • \$1,700,000		XTX Ventures Undisclosed
	Illuminate Financial Management Undisclosed		

All documents

Q Search

Download all

Document	Uploaded	
 Disclaimers.pdf	Apr 3, 2026	
 Closing documents	Apr 4, 2026	
 Gladia Deck 2026	Apr 3, 2026	