

Building the compute economy

Mike Henry, CEO & Founder
Tim Harris, Co-founder

proprietary & confidential



Shift to a Compute-Driven AI Economy



Demand for AI is causing a global rebuild of compute infrastructure



Consumption is also changing: AI developers consume compute in the form of tokens and higher levels of AI work




The scale of compute needed dwarfs what legacy cloud vendors can provide, leading to a vast landscape of new AI cloud factories



These AI cloud factories are largely siloed and cannot bridge the gap between compute and AI work



Parasail has **unified a fragmented landscape of AI factories** powering the fastest growing AI companies with **high performance, low-cost AI inference**



Parasail Inference Platform

Demand layer



Customer



Customer



Customer

Parasail is an
unparalleled
AI Inference
Platform

Supply layer



Alternative AI
HW Clouds



Neo Cloud



Neo Cloud



Alternative AI HW
Clouds

Built on an
AI Deployment
Network of the
world's best
compute clouds

Parasail Inference Platform



**Parasail is an
unparalleled
AI Inference
Platform**



Customer



Units of AI Work



Parasail Inference Platform

Performance

Geography

Latency

Hardware
Selection

Price

Security and
Compliance

Parasail Inference Platform

Demand: The Best Inference Platform

- 75B tokens a day, growing fast
- Zero churn of production customers
- Rapid growth in 5 months public
 - \$3M ARR in 5 month
 - 30 enterprise customers, thousands of developers

Parasail is an
unparalleled
AI Inference
Platform

Customer

Use Cases of AI Work

Parasail Inference Platform

Geography

Latency


Selection

Price

Security and
Compliance

Parasail Inference Platform

Planetary Kubernetes



Supply layer

Alternative AI
HW Clouds

Neo Cloud

Neo Cloud

Alternative AI HW
Clouds

Built on an
**AI Deployment
Network** of the
world's best
compute clouds

Parasail Inference Platform

Planetary Ku

Supply-Side Maximalism

- 25 compute providers, growing fast
- Enough capacity potential to overnight build a GPU fleet bigger than Oracle
- 50% gross margin, when competitors are struggling to be margin positive

Built on an
**AI Deployment
Network** of the
world's best
compute clouds

Exponential Growth for Inference Compute

Cloud Compute Market

\$1.5T

By 2032



Inference CAGR

~19.2%

\$365B by 2032




Open-source Usage

30%→70%

Adoption by 2032



Inference market projected to reach \$365bn by 2032



Performance with Flexibility, Scale without the Cost

Demand-side Inference Platform

Flexible

Designed to flexibly deploy specialized AI workflows

Performance

Synthetic scale and diversity drive most parameters for optimization


Supply-side AI Deployment Network

Cost


Competition from 100+ data centers and clouds drives best price

Scale

Planetary scale Kubernetes gives us access to fleet of on-demand GPUs larger than Oracle's



The Standard LLM Inference Stack



Why did we feel the need to build something different?

Maximal stance on supply side is key to winning inference.

Managing Kubernetes is hard: if you have 10 neoclouds and 10 regions each, that's 100 clusters to manage.

What if a neocloud doesn't have Kubernetes?

What if a neocloud is unreliable?

Planetary Kubernetes

Orchestrating the Earth's GPU resources from a single plane




Scale: global hyper-scale without buying any hardware



Reach: access to lesser-known datacenters with low utilization




Diversity: new NVIDIA, old NVIDIA, new AI HW platforms




The Parasail Deployment Network

Control Plane (AWS + Azure)



Compute Plane



Planetary Kubernetes

- Unified DevOps and security operations across many datacenters and providers
- Unparalleled utilization and efficiency
- Easy to work with tier 2 vendors
- Multicloud that is invisible to the customer

Building the Best Inference Platform

Deployment Network Foundation

Scale

Synthetic scale provides optimization parameters and economies of Scale


Reach

Time, Geography, Providers

Diversity

The latest NVIDIA, new AI platforms, deprecated hardware, etc.

Automatic Optimization



Building the Best Inference Platform

Multi-Engine Inference

Optimal Engine


- Match the best inference engine for the right model and application


Leverage Open Source

- Open source is outpacing private development

Invisible to Customer

- Seamless to the customer, simply delivers the best performance





Building the Best Inference Platform

Model Optimization

Tuned to Application


- Right model (open or proprietary), tuned makes inference sticky with customer

Design Automation

- Use automation and lots of compute to search design space

Kernel Tuning

- Only bring in kernel tuning on a case-by-case basis



Building the Best Inference Platform

Dataflow Processing

Dynamic Batching


- Weaving of batch process in GPU downtime to achieve huge savings by optimizing usage

Stream Processing for Agents

- Time threading of agent flows for seamless experience

Complete Utilization

- Achieves near-perfect fleet utilization, in industry with 30-50% average idle time



Building the Best Inference Platform

Agentic Workflows

Search and Research


Interactive and reflective queries, aka “research”

Stateless Agent Deployment

Serverless hosting of stateless orchestrators

Browser Control

Autonomous navigation of the web



Building the Best Inference Platform

Complex Modalities

Web Search


- Up to date web search and scraping

Generative Media

- RAG pipelines for generating media

Voice Pipeline

- Real time high performance voice through proprietary model partnerships





Met **last Monday**, production ramp **last Friday**,
\$500k ARR by EOY

POSITRON

\$5M LOI to manage 42 racks of hardware
starting in February

ACCELERATION



Special products on their AI gateway

In the last
two weeks...



Winning their production ramp of Hermes 4
over **Lamba Labs** and **Nebius**



Hosting Allen Institute Olmo models
+ collaborating on internal projects

Sandbar

Closed yesterday, migrating voice
traffic from Groq to us, 10k user in Q1 2026

Traction

Public Launch April 2025

- Thousands of developers
- 25 enterprise customers
- 10 day average to sign
- Fleet of on-demand GPU's larger than Oracle

Launch April

Today

\$3MM
ARR

Year End

\$8MM

ARR 2025E

\$5M

from existing
customers

April

August

December

Agentic Memory with No Limits to Scale

Humanlike Multimodal
Agent Company

Challenge

- Could not get the rate limit they needed from OpenAI. Decided to migrate all traffic to GPT-OSS.

Solution

- Rapid onboarding, after which optimization immediately kicked in

Results


- First meeting Monday, running production traffic on Friday.

Rate Limits

Eliminated

\$500k

ARR by EOY



Unlocking a Flagship Deal with Lightning-Fast Inference Deployment



Challenge

- High inference costs and operational complexity blocked a major deal with a top food delivery service

Solution

- Delivered a fully optimized, hosted inference deployment in just 10 days without compromise on quality, eliminating internal engineering burden

Results


- Unlocked flagship deal and \$250K savings, establishing Parasail as the preferred infrastructure partner for 2025 expansion

17

Days to contract

\$500k

ARR by EOY



Scaling Elicit's Scientific Copilot for 2M+ Researchers



Challenge

- ▶ Limited product control, high costs, and inflexible infrastructure due to reliance on OpenAI

Solution

- ▶ Engaged Parasail to gain flexible, cost-efficient GPU access with simplified orchestration and integration

Results


- ▶ Lower latency and costs while freeing engineering teams to focus on core R&D innovation

25B

Tokens Served Daily

\$2M

ARR by EOY



Led by repeat founders with the technical depth and commercial intuition to win

leadership

Former founder and CEO of Mythic, raised \$165M to build a disruptive AI inference compute platform. Passionate innovator who transforms ideas into impact.



Mike Henry
CEO &
Co-Founder



Tim Harris
Co-Founder &
Board Director

Founder and CEO of Swift Navigation. Expert deal-maker and business development leader. Built teams and products driving hundreds of millions in sales and raised \$250M in venture capital.



team

6 backend engineers, veterans of...



supply-side

4 performance engineers, veterans of...



demand-side

Use of Fundraise

Focus on Platform

- Scaling engineering across AI and Platform, with emphasis on workflows

Scale GTM

- Build first GTM team (sales, marketing, business dev)

Key hires

- Head of Business Development (Supply Side), Head of Marketing, Head of Sales (Demand Side)

3 → 12

GTM

10 → 25

Engineering

ARR Projection through 2027



EOY 2026 target of \$28M ARR with EOY
2027 target of \$67M ARR